

Parameter estimation for the gestalt model

1 Generative model

A gestalt, a perceptual object, is characterised by a covariance component for the joint distribution of visual neural activity.

$$p(v \mid g) = \mathcal{N}(v; 0, C_v) \quad (1)$$

$$C_v = \sum_{k=1}^K g_k C_k \quad (2)$$

where K is the fixed number of possible gestalts in the visual scene and g_k is the strength of the gestalt number k , coming from a K -dimensional symmetric Dirichlet prior distribution with concentration parameter α controlling the sparsity of the prior.

$$p(g) = \text{Dir}(g; \alpha) \quad (3)$$

The pixel intensities are generated from the neural activity through a set of linear projective field models, possibly Gabor filters, A , adding some independent observational noise.

$$p(x \mid v) = \mathcal{N}(x; Av, C_x) \quad (4)$$

$$C_x = \sigma_x I; \quad (5)$$

We might assume that a single composition of gestalts, characterised by the g vector, generates a batch of B independent images, described by cellular activities $V = \{v_1 \dots v_B\}$ and observations $X = \{x_1 \dots x_B\}$. Then the following likelihood distributions hold

$$p(V \mid g) = \prod_{b=1}^B \mathcal{N}(v_b; 0, C_v) \quad (6)$$

$$p(X \mid V) = \prod_{b=1}^B \mathcal{N}(x_b; Av_b, C_x) \quad (7)$$

The likelihood of the covariance components with respect to a dataset of size N , $\mathbf{X} = \{X_1 \dots X_N\}$ can be expressed as follows

$$\begin{aligned}
p(\mathbf{X} \mid C_1 \dots C_k) &= \int \int_{-\infty}^{\infty} p(\mathbf{X} \mid \mathbf{V}) p(\mathbf{V} \mid G) p(G) d\mathbf{V} dG = \\
&= \int_{-\infty}^{\infty} \prod_{n=1}^N p(g_n) \int_{-\infty}^{\infty} \prod_{b=1}^B p(x_{n,b} \mid v_{n,b}) p(v_{n,b} \mid g_n) d\mathbf{V} dG
\end{aligned} \tag{8}$$

2 Gibbs sampling as the E-step

An efficient way to collect samples from the joint posterior over all hidden variables is to employ a Gibbs sampling scheme, where we sample from the conditional posteriors. The first is over v , and can be defined as follows

$$p(v \mid x, g) = \frac{p(x \mid v, g) p(v \mid g)}{p(x \mid g)} = \frac{\mathcal{N}(x; Av, \sigma_x I) \mathcal{N}(v; 0, C_v)}{\int_{-\infty}^{\infty} \mathcal{N}(x; Av, \sigma_x I) \mathcal{N}(v; 0, C_v) dv} \tag{9}$$

The Gaussian over x can be rewritten to a Gaussian over v times a constant c_1 in the following way

$$\mathcal{N}(x; Av, \sigma_x I) = c_1 \mathcal{N}(v; (A^T A)^{-1} A^T x, \sigma_x (A^T A)^{-1}) = c_1 \mathcal{N}(v; A^+ x, \sigma_x (A^T A)^{-1}) \tag{10}$$

where A^+ is the Moore-Penrose pseudoinverse of A , with $D_x = D_v \rightarrow A^+ = A^{-1}$. Consequently, the product of two Gaussians in the numerator of Eq. 8 can also be written as a Gaussian over v introducing a new constant

$$\mathcal{N}(x; Av, \sigma_x I) \mathcal{N}(v; 0, C_v) = c_1 c_2 \mathcal{N}(v; \mu_{post}, C_{post}) \tag{11}$$

The denominator of Eq. 8 is the integral of this formula, which evaluates to $c_1 c_2$, as the Gaussian integrates to one. This cancels the constants in the numerator, making the conditional posterior equal to the combined Gaussian over v , which, after expanding μ_{post} and C_{post} , is

$$p(v \mid x, g) = \mathcal{N} \left(v; \frac{1}{\sigma_x} \left(\frac{1}{\sigma_x} A^T A + C_v^{-1} \right)^{-1} A^T x, \left(\frac{1}{\sigma_x} A^T A + C_v^{-1} \right)^{-1} \right) \tag{12}$$

and for a batch of size B

$$p(V \mid X, g) = \prod_{b=1}^B \mathcal{N}(v_b; \mu_{post}(x_b), C_{post}) \tag{13}$$

which can be sampled directly from a Gaussian of dimension $D_v \times B$. The conditional posterior over g is defined as follows

$$p(g \mid X, V) = \frac{p(X \mid g, V) p(g \mid V)}{p(X \mid V)} = \frac{p(V \mid g) p(g)}{p(V)} \tag{14}$$

which can be sampled by a slice sampling scheme with the following target

$$\log p(g \mid X, V) \sim -\frac{1}{2} \left[B \log(\det(C_v)) + \sum_{b=1}^B v_b^T C_v^{-1} v_b \right] + (\alpha - 1) \sum_{k=1}^K \log(g_k) \quad (15)$$

3 M-step

The complete-data likelihood with respect to a set of batch observations of size B is the following

$$\begin{aligned} p(\mathbf{V}, G, \mathbf{X} \mid C_{1..K}) &= \prod_{n=1}^N p(X_n \mid V_n) p(V_n \mid g_n) p(g_n) = \\ &= \prod_{n=1}^N p(g_n) \prod_{b=1}^B p(x_{nb} \mid v_{nb}) p(v_{nb} \mid g_n) \end{aligned} \quad (16)$$

Let's denote the logarithm of this by $\mathcal{L} = \log p(\mathbf{V}, G, \mathbf{X} \mid C_{1..K})$. We can approximate the integral of this logarithm over the joint posterior by averaging over L samples from it, separately for each observation x_n . As we will seek the values of the precision components $C_{1..K}$ that maximise this integral, we can discard each term not depending on these parameters. This way we arrive to the following expression

$$\begin{aligned} \mathcal{L} &\sim \sum_{n=1}^N \frac{1}{L} \sum_{l=1}^L -\frac{1}{2} \left[B \log \left(\det \left(C_v^{(l,n)} \right) \right) + \sum_{b=1}^B v^{(l,n,b)T} \left(C_v^{(l,n)} \right)^{-1} v^{l,n,b} \right] = \\ &= -\frac{1}{2L} \sum_{m=1}^{NL} \left[B \log \left(\det \left(C_v^m \right) \right) + \sum_{b=1}^B v^{(m,b)T} \left(C_v^m \right)^{-1} v^{m,b} \right] \end{aligned} \quad (17)$$

noting that the double summation over L samples over all N observations always happens on the same terms, so we can substitute it with a single sum that iterates over the full sample set.

To ensure that the optimisation procedure does not produce precision matrices that are not positive definite, we can optimise for the Cholesky upper triangle matrix instead of the precision matrix, as this also specifies the Gaussian completely.

$$C_j = U_j^T U_j \quad (18)$$

$$C_v = \sum_{k=1}^K g_k U_k^T U_k \quad (19)$$

$$\frac{\partial C_v^m}{\partial U_j} = 2g_j^m U_j \quad (20)$$

so by the chain rule, the derivative of \mathcal{L} according to U_j looks like this

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial U_j} &= -\frac{1}{2L} \sum_{m=1}^{NL} \frac{\partial \mathcal{L}^m}{\partial C_v^m} \frac{\partial C_v^m}{\partial U_j} = \\ &= -\frac{1}{L} \sum_{m=1}^{LN} g_j^m \left[B (C_v^m)^{-1} - \sum_{b=1}^B (C_v^m)^{-1} v^{m,b} v^{(m,b)T} (C_v^m)^{-1} \right] U_j = \\ &= -\frac{1}{L} \sum_{m=1}^{LN} g_j^m \left[B (C_v^m)^{-1} - (C_v^m)^{-1} \left(\sum_{b=1}^B v^{m,b} v^{(m,b)T} \right) (C_v^m)^{-1} \right] U_j \end{aligned} \quad (21)$$

As a generalised M-step, we can move the parameters in the direction of the gradient scaled by a learning rate

$$U_j^{new} = U_j^{old} + \epsilon \frac{\partial \mathcal{L}}{\partial U_j} \quad (22)$$

4 Parametrisation with precision components

The model can be equally well parametrised by precision components

$$p(v | g) = \mathcal{N}(v; 0, \Lambda_v^{-1}) \quad (23)$$

$$\Lambda_v = \sum_{k=1}^K g_k \Lambda_k \quad (24)$$

in this case the conditional posterior over v takes the form

$$p(v | x, g) = \mathcal{N} \left(v; \frac{1}{\sigma_x} \left(\frac{1}{\sigma_x} A^T A + \Lambda_v \right)^{-1} A^T x, \left(\frac{1}{\sigma_x} A^T A + \Lambda_v \right)^{-1} \right) \quad (25)$$

and the sampling target for the conditional posterior of g will look as follows

$$\log p(g | X, V) \sim -\frac{1}{2} \left[B \log(\det(\Lambda_v^{-1})) + \sum_{b=1}^B v_b^T \Lambda_v v_b \right] + (\alpha - 1) \sum_{k=1}^K \log(g_k) \quad (26)$$

The gradient of the expectation of the complete-data log-likelihood with respect to the joint posterior will look like this

$$\Lambda_j = U_j^T U_j \quad (27)$$

$$\frac{\partial \mathcal{L}}{\partial U_j} = \frac{1}{L} \sum_{m=1}^{LN} g_j^m \left[B (\Lambda_v^m)^{-1} - \sum_{b=1}^B v^{m,b} v^{(m,b)T} \right] U_j \quad (28)$$